Topic 3 – Discrete Distributions

ENVX1002 Introduction to Statistical Methods

Dr. Floris van Ogtrop
The University of Sydney

Jan 2024

THE UNIVERSITY OF SYDNEY

Outline – Discrete distributions

- Example

- What is a distribution

- Binomial distribution

- Poisson distribution

Learning outcomes

At the end of this topic students should able to:

- Have a good understanding of what a distribution is

  - Definitions

  - Functions

  - Binomial and Poisson distributions

- Apply the correct model to describe data

- Demonstrate proficiency in the use of R and Excel for calculating probabilities

THE UNIVERSITY OF SYDNEY

Types of data

Remember the types of data

Numerical

- Continuous: yield, weight

- Discrete: weeds per m^2

Categorical

- Binary: 2 mutually exclusive categories

- Ordinal: categories ranked in order

- Nominal: qualitative data

THE UNIVERSITY OF SYDNEY

Example

- We have 5 insects which we spray with an insecticide, each insect has a 60% chance of being killed;

- $P(K) = 0.6$

- Possible questions:
  - What is the probability that all 5 insects will be killed?
  - What is the probability that at least 3 insects will be killed?

- The data is 'binary' and the events are mutually exclusive (either dead or alive unless it is zombie fly);
  - we can say the data is categorical or numeric discrete
  - we can use a binomial (discrete) distribution to "model" the data.

Generated using DALL.E3

THE UNIVERSITY OF SYDNEY

What is a distribution?

- In our case we are generally referring to a distribution function

  - This is a function (or model) that describes the probability that a system will take on value or set of values {x}

- For any variable X, we describe probabilities by

  - Discrete variables: probability distribution function $P(X = x)$

  - Continuous variables: probability density function $f(x)$

  - Discrete and Continuous variables: cumulative density function $F(x) = P(X \leq x)$

Back to our example

- We spray 5 flies with insecticide which has 60% chance of killing each insect. If X is the number of flies that die, what is the distribution of X?

- The set of possible values is $x = 0, 1, 2, 3, 4, 5$

- The likelihood of each value is

- $P(X = 0) = P(\text{no insects die}) = 0.4 \times 0.4 \times 0.4 \times 0.4 \times 0.4 = 0.4^5 = 0.01024$

- $P(X = 1) = P(\text{one insects die}) = 0.6 \times 0.4 \times 0.4 \times 0.4 \times 0.4 + 0.4 \times 0.6 \times 0.4 \times 0.4 \times 0.4 + 0.4 \times 0.4 \times 0.6 \times 0.4 \times 0.4 + 0.4 \times 0.4 \times 0.4 \times 0.6 \times 0.4 + 0.4 \times 0.4 \times 0.4 \times 0.4 \times 0.6 = 0.0768$

- $P(X = 2) = P(\text{two insects die}) = \ldots 10 \text{ different combinations} = 0.2304$

- $P(X = 3) = P(\text{three insects die}) = \ldots 10 \text{ different combinations} = 0.3456$

- $P(X = 4) = P(\text{four insects die}) = \ldots 5 \text{ combinations} = 0.2592$

- $P(X = 0) = P(\text{all insects die}) = 0.6 \times 0.6 \times 0.6 \times 0.6 \times 0.6 = 0.65 = 0.07776$

Example plot

```
1  x <- c(0, 1, 2, 3, 4, 5)
2  p <- c(0.01024, 0.0768, 0.2304, 0.3456, 0.2592, 0.07776)
3  plot(x, p, type = "h")
```

Example – properties of the distribution

- Remember we have a binomial (dead or alive) distribution here

- A key property of all (discrete) distributions is that all probabilities add to

$$\Sigma_{i=0}^{5} P(X = i) = 0.01024 + 0.0768 + 0.2304 + 0.3456 + 0.2592 + 0.07776 = 1$$

- Note that all probabilities lie between 0 and 1

THE UNIVERSITY OF SYDNEY

Binomial Distribution

- Why does a binomial distribution fit our insect data??

    - Basic element is a Bernoulli trial – each insect;

    - The outcome of each trial can be classified in precisely one of two mutually exclusive ways termed "success" (dead) and "failure" (alive);

        ○ We usually assign p to success and q to failure.

    - Binomial experiments consists of n Bernouilli (independent binary) trials (i.e. 5 insects);

    - The probability of a success, denoted by p, remains constant from trial to trial. The probability of a failure, $q = 1-p$;

        ○ $p = 0.6$ and $q = 0.4$

    - The trials are independent; that is, the outcome of any particular trial is not affected by the outcome of any other trial;

    - The number of successes, x, is a binomial variable.

THE UNIVERSITY OF SYDNEY

Example

- How many combinations are there for exactly 2 flies to die out of 5 flies?

$$\binom{5}{2} = \frac{5!}{2!(5-2)!} = \frac{(5 \times 4 \times 3 \times 2 \times 1)}{(2 \times 1 \times 3 \times 2 \times 1)} = 10$$

- What is the probability that exactly 2 flies will die?

$$P(X = x) = \binom{n}{x} p^x (1-p)^{n-x} = \binom{5}{2} 0.6^2 (1 - 0.6)^{5-2}$$

$$= 10 \times 0.36 \times 0.064 = 0.2304$$

- `dbinom(2,5,0.6)`

```
1  dbinom(2,5,0.6)
```

```
[1] 0.2304
```

- `=BINOM.DIST(2,5,0.6,FALSE)`

Point, cumulative or interval probabilities

- We have already calculated a point probability i.e. the probability of exactly 2 flies dying.

- But what about if we wanted to know the probability of 2 or more flies dying $P(X \geq 2)$ or between 2 and 4 flies $P(2 \leq X \leq 4)$ dying??

- So how can we calculate these?
  - One way is to calculate all the point probabilities from 0-5 and add the probabilities cumulatively or in the interval

x	P
0	0.01024
1	0.0768
2	0.2304
3	0.3456
4	0.2592
5	0.07776
SUM	1

Point, cumulative or interval probabilities

$P(X \geq 2) = (X = 2) + (X = 3) + (X = 4) + (X = 5) = 0.2304 + 0.3456 + 0.2592 + 0.07776 = 0.91296$

- You guys try to calculate the following

$P(2 \leq X \leq 4) = ??$

x	P
0	0.01024
1	0.0768
2	0.2304
3	0.3456
4	0.2592
5	0.07776
SUM	1

THE UNIVERSITY OF SYDNEY

Cumulative

$P(X \geq 2)$


```
1  1 - pbinom(1,5,0.6)
```

```
[1] 0.91296
```

```
1  ## OR
2
3  pbinom(1,5,0.6, lower.tail = FALSE)
```

```
[1] 0.91296
```


`=1-BINOM.DIST(1,5,0.6,TRUE)`

Interval

$P(2 \leq X \leq 4)$


```
1  pbinom(4,5,0.6)-pbinom(1,5,0.6)
```

```
[1] 0.8352
```



```
=BINOM.DIST(4,5,0.6,TRUE)-BINOM.DIST(1,5,0.6,TRUE)
```

Mean and variance of the binomial distribution

- Mean binomial distribution

$$\mu_x = np$$

$= 5 \times 0.6 = 3$ On average 3 flies die in 5 trials

- Variance binomial distribution

$$\sigma_x^2 = np(1-p)$$

$= 5 \times 0.6(1 - 0.6) = 1.2$ with a variance of 1.2 flies

Count Data

- Often we are interested in count data such as the number of events occurring in an interval;

- We generally model this data using the Poisson distribution (Described by Simeon-Denis Poisson) where we use λ denotes the average number of events occurring in an interval

- We often write this as

$$X \sim Po(\lambda)$$

- Where X is the number of discrete and independent events in the interval

  - e.g. number of plants of a certain species along transect

  - e.g. occurrence of disease in a period of time

Horse kick deaths in the Prussian Army

```
1  library(knitr)
2  kick <- read.csv("data/Kick_deaths.csv")
3  kable(kick[1:12,])
```

Year	GC	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	C11	C14	C15
1875	0	0	0	0	0	0	0	1	1	0	0	0	1	0
1876	2	0	0	0	1	0	0	0	0	0	0	0	1	1
1877	2	0	0	0	0	0	1	1	0	0	1	0	2	0
1878	1	2	2	1	1	0	0	0	0	0	1	0	1	0
1879	0	0	0	1	1	2	2	0	1	0	0	2	1	0
1880	0	3	2	1	1	1	0	0	0	2	1	4	3	0
1881	1	0	0	2	1	0	0	1	0	1	0	0	0	0
1882	1	2	0	0	0	0	1	0	1	1	2	1	4	1
1883	0	0	1	2	0	1	2	1	0	1	0	3	0	0
1884	3	0	1	0	0	0	0	1	0	0	2	0	1	1
1885	0	0	0	0	0	0	1	0	0	2	0	1	0	1
1886	2	1	0	0	1	1	1	0	0	1	0	1	3	0

THE UNIVERSITY OF SYDNEY

Horse kick deaths in the Prussian Army

https://en.wikipedia.org/wiki/Ladislaus_Bortkiewicz

```
1  library(tidyverse)
2
3  frequency_kick <- kick %>%
4    select(-Year) %>%
5    pivot_longer(cols = everything(), names_to = "Column",
6    count(Deaths) %>%
7    arrange(Deaths) %>%
8    mutate(Total_Deaths = Deaths*n) %>%
9    mutate(Probability = "?")
10
11 kable(frequency_kick)
```

Deaths	n	Total_Deaths	Probability
0	144	0	?
1	91	91	?
2	32	64	?
3	11	33	?
4	2	8	?

- What is the Probability in any month of
  - 0 injuries by horse kick
  - 1 injuries by horse kick
  - 2 injuries by horse kick
  - 3 injuries by horse kick
  - 4 injuries by horse kick
- λ "Lambda" is the mean

THE UNIVERSITY OF SYDNEY

Horse kick deaths in the Prussian Army

```
1  total_kick <- frequency_kick %>%
2    summarize(n = sum(n), sum_Total_Deaths = sum(Total_Deaths))
3
4  kable(total_kick)
```

n	sum_Total_Deaths
280	196

Poisson Distribution

$X \sim Po(\lambda)$

$P(X = x) = \frac{\lambda^x e^{-\lambda}}{x!} \quad x = 0, 1, 2, \ldots \lambda > 0$

- Note that e denotes the exponential function such that

  - $e^0 = 1$

  - $e^-2 = 0.135 (3 \text{ d.p.})$

  - $e^-10 = 4.540 \times 10^{-5} (3 \text{ d.p.})$

THE UNIVERSITY OF SYDNEY

Example with kick deaths

- We first identify the model

  - X = the number of soldiers injured by horse kick $\sim Po(\lambda)$ where λ = the average number of deaths = $196/280 = 0.7$. We can now calculate the probability of having exactly 0, 1, 2, 3, 4, 5 deaths

- $P(X=0) = \frac{0.7^0 e^{-0.7}}{0!} = \frac{1e^{-0.7}}{1} = 0.497 (3 \text{ d.p.})$
- $P(X=1) = \frac{0.7^1 e^{-0.7}}{1!} = \frac{0.7 e^{-0.7}}{1} = 0.348 (3 \text{ d.p.})$
- $P(X=2) = \frac{0.7^2 e^{-0.7}}{2!} = \frac{0.49 e^{-0.7}}{2 \times 1} = 0.122 (3 \text{ d.p.})$
- $P(X=3) = \frac{0.7^3 e^{-0.7}}{3!} = \frac{0.343 e^{-0.7}}{3 \times 2 \times 1} = 0.028 (3 \text{ d.p.})$
- $P(X=4) = \frac{0.7^4 e^{-0.7}}{4!} = \frac{0.2401 e^{-0.7}}{4 \times 3 \times 2 \times 1} = 0.005 (3 \text{ d.p.})$

Example with kick deaths

```
1  x <- c(0, 1, 2, 3, 4)
2  p <- c(0.497, 0.348, 0.122, 0.028, 0.005)
3  plot(x, p, type = "h")
```

Example with kick deaths

```
 1  frequency_kick1 <- kick %>%
 2    select(-Year) %>%
 3    pivot_longer(cols = everything(), names_to = "Column", values_to = "Deaths") %>%
 4    count(Deaths) %>%
 5    arrange(Deaths) %>%
 6    mutate(Total_Deaths = Deaths*n) %>%
 7    mutate(Probability = c(0.497, 0.348, 0.122, 0.028, 0.005)) %>%
 8    mutate(Observed_Probability = n/280)
 9
10  kable(frequency_kick1)
```

Deaths	n	Total_Deaths	Probability	Observed_Probability
0	144	0	0.497	0.5142857
1	91	91	0.348	0.3250000
2	32	64	0.122	0.1142857
3	11	33	0.028	0.0392857
4	2	8	0.005	0.0071429

- note that observed probability is n divided by the total number of observations (280). For example, for 0 deaths there were 144 observed out of a total of 280 observations i.e. the observed probability of 0 deaths in a cavalry corps over 20 years of observations was 0.51 or 51%

THE UNIVERSITY OF SYDNEY

Example with kick deaths

- So now you all can calculate what the probability is, as an example, the of having less than 2 deaths across all cavalry corps for the period of 1875-1894 $P(X < 2)$.


```
1  ppois(1,0.7)
```

```
[1] 0.844195
```



```
=POISSON(1,0.7,TRUE)
```

Example with kick deaths

- Another example, is having exactly 2 deaths across all cavalry corps for the period of 1875-1894 $P(X = 2)$.

```
1  dpois(1,0.7)
```

```
[1] 0.3476097
```

```
- =POISSON(2,0.7,FALSE)
```

Interesting results with the binomial and Poisson distributions

- For large n and small p the Binomial distribution $X \sim Bin(n, p)$ can be approximated by the Poisson distribution $Y \sim Po(\lambda = np)$

  - The general rule is if $n > 20$ and $np < 5$

- We often say that the Poisson Distribution models rare events

Interesting results with the binomial and Poisson distributions

```r
1  # Parameters for the binomial distribution
2  n <- 100
3  p <- 0.05
4
5  # Calculating lambda for the Poisson approximation
6  lambda <- n * p
7
8  # Generate the range of values
9  x <- 0:n
10
11 # Data frames for plotting
12 data_binom <- data.frame(x = x, probability = dbinom(x, n
13 data_pois <- data.frame(x = x, probability = dpois(x, laml
14
15 # Combine data
16 data_combined <- rbind(data_binom, data_pois)
17
18 # Create the plot
19 p <- ggplot(data_combined, aes(x = factor(x), y = probabil
20     geom_bar(stat = "identity", position = position_dodge
21     ggtitle("Binomial vs Poisson Distribution") +
22     xlab("Number of successes") +
23     ylab("Probability") +
24     scale_fill_manual(values = c("blue", "red")) +
```

```r
1  print(p)
```


Binomial vs Poisson Distribution

Further reading

- Quinn & Keough (2002)
    - Chapter 1. Sections 1.5, p. 9-13
- Mead et al. (2002)
    - Chapter14. Sections 14.4-14.5, p. 339-377

THE UNIVERSITY OF SYDNEY

Thanks!

This presentation is based on the SOLES Quarto reveal.js template and is licensed under a Creative Commons Attribution 4.0 International License.

THE UNIVERSITY OF SYDNEY